

## Machine Problem 4 Report

Lab members: Martin Fracker, Chris Findeisen

Date: April 18, 2015

## **Introduction**

During machine problem 4, we worked with threads to extend our dataserver/client relationship from machine problem 2. We first forked off a process to load the data server. We then called a client that used three request threads and worker threads. These threads used a queue mutex and two semaphores to access a bounded buffer. Finally, we joined the threads when their work was done.

## **Procedures**

We implemented the three different types of threads, and implemented bounded buffer. We created these in our [main.cc](#) and joined them at the end before sending “quit” to the data server.

## Results:

We used ascii-engine, built by Martin, to display our results, which allowed us to update in place and in real time. Below are two snapshots of it. The first is while the 10,000 requests was running and afterwards, when it was finished.

Press q to exit display

Range:	Joe Smith		Jane Smith		John Doe
1-10	307		289		303
11-20	282		313		303
21-30	279		286		270
31-40	276		296		304
41-50	264		300		300
51-60	328		293		324
61-70	304		316		288
71-80	312		299		310
81-90	319		287		270
91-100	300		321		297
Total	2971		3000		2969

Press q to exit display

Range:	Joe Smith		Jane Smith		John Doe
1-10	1002		976		970
11-20	967		996		1017
21-30	991		983		959
31-40	980		993		980
41-50	907		998		974
51-60	1017		1048		1066
61-70	954		1052		947
71-80	1094		994		1039
81-90	1037		975		1005
91-100	1051		985		1043
Total	10000		10000		10000

## Performance Analysis:

The client consists of the worker threads, client threads, and the histogram threads, used for statistics.

Worker threads dequeue and client threads enqueue. Context switching when there are more threads than cores, can cause performance decreases. However, because only one thread can access the queue at one time (because of the mutex), only one thread can run at a single time during that crucial section. We identify this as the biggest performance loss, since worker threads and client threads mostly run on these instructions.

```
$ python performance_tester.py
----- buffer size = 10 -----
worker thread count = 5
5.11588096619
worker thread count = 10
2.5437669754
worker thread count = 25
1.02014303207
worker thread count = 1000
0.543351888657
----- buffer size = 100 -----
worker thread count = 5
5.09391403198
worker thread count = 10
2.52165794373
worker thread count = 25
1.01570892334
worker thread count = 1000
0.541793107986
----- buffer size = 1000 -----
worker thread count = 5
5.00269412994
worker thread count = 10
2.46230697632
worker thread count = 25
0.997351884842
worker thread count = 1000
0.5762591362
----- buffer size = 10000 -----
worker thread count = 5
4.91797208786
worker thread count = 10
2.45690703392
worker thread count = 25
0.996287822723
worker thread count = 1000
0.508797883987
```

<- This figure shows the performance of our program with differing buffer and worker sizes. The times are shown in seconds. As you can see, the worker thread count greatly increased efficiency. We found that when the worker thread count went over 1000, the efficiency returns began to be almost nil. The increased buffer size also helped, although less dramatically than increasing the worker thread count.

## **Conclusion**

We learned a lot about threading, especially debugging concurrency issues. One take-away for us was to carefully pair-program over critical sections in our code, in order to catch silly mistakes. Likewise, when there is a problem, the method to handle it should involve 90% reading, and 10% coding.

If we were to do this again, we would have split the threads up less granularly, giving them a longer sequence of instructions. This would allow our code to benefit more greatly from our threading.